

```
import pandas as pd

from sklearn.metrics import mean_squared_error, r2_score

from sklearn.model_selection import train_test_split, RepeatedStratifiedKFold, c

from sklearn.linear_model import LinearRegression

import numpy as np

import re

import matplotlib.pyplot as plt

from sklearn.linear_model import Lasso

from sklearn.linear_model import Ridge

from sklearn.ensemble import RandomForestRegressor

import seaborn as sns

# read into pandas dataframe
df = pd.read_csv("used_car_sales.csv")
```

df

	ID	pricesold	yearsold	zipcode	Mileage	Make	Model	Year
0	137178	7500	2020	786**	84430	Ford	Mustang	1988
1	96705	15000	2019	81006	0	Replica/Kit Makes	Jaguar Beck Lister	1958
2	119660	8750	2020	33449	55000	Jaguar	XJS	1995
3	80773	11600	2019	07852	97200	Ford	Mustang	1968
4	64287	44000	2019	07728	40703	Porsche	911	2002
...
122139	14948	4200	2019	80233	102700	Ford	Mustang	1977
122140	58814	6500	2019	53132	128000	Ford	E-Series Van	2012
122141	2156	2000	2019	77536	50000	Ford	Bronco	1978
122142	29096	2280	2019	92131	164337	BMW	3-Series	2000
122143	52391	5000	2019	18951	163111	Nissan	300ZX	1990


```
df.describe()
```

	ID	pricesold	yearsold	Mileage	Year	Num
count	122144.000000	122144.000000	122144.000000	1.221440e+05	1.221440e+05	1.2
mean	85094.212397	10808.560715	2019.375467	1.404291e+06	3.959362e+03	1.7
std	47786.970812	13987.295760	0.503671	3.335593e+07	1.984514e+05	6.7
min	1.000000	0.000000	2018.000000	0.000000e+00	0.000000e+00	0.0
25%	44547.250000	2950.000000	2019.000000	4.479225e+04	1.977000e+03	4.0
50%	85555.500000	6500.000000	2019.000000	9.000000e+04	2.000000e+03	6.0
75%	127078.500000	13800.000000	2020.000000	1.402000e+05	2.008000e+03	8.0
max	165801.000000	404990.000000	2020.000000	1.235669e+09	2.014000e+07	2.7

```
# Plotting the histogram with more bins
plt.hist(df['pricesold'], bins=90, color='blue', edgecolor='black')

# Adding labels and title
plt.xlabel('Price Sold')
plt.ylabel('Frequency')
plt.title('Histogram of Price Sold')


# Displaying the plot
plt.show()
```



```
# Plotting the histogram with more bins
plt.hist(df['yearsold'], bins=90, color='blue', edgecolor='black')

# Adding labels and title
plt.xlabel('Vehicle Age')
plt.ylabel('Density')
plt.title('The Density Histogram of Vehicle Age')

# Displaying the plot
plt.show()
```



```
# read into pandas dataframe
df = pd.read_csv("used_car_sales.csv")

# dropping
to_drop = ['ID', 'zipcode', 'Trim']

df.drop(to_drop, inplace=True, axis=1)
```

```
# types left
print(df.dtypes)
# pricesold, yearsold, mileage, make, model, year

# drop missing values
#df.dropna()

count_initial = len(df.index)
print(count_initial)

# max and min price
max_price = 50000
min_price = 1000
df = df[max_price > df.pricesold]
df = df[min_price < df.pricesold]
count_price = len(df.index)
print("price removed: " + str(count_initial - count_price) + " outliers")

# max and min mileage
max_mileage = 200000
min_mileage = 1000
df = df[max_mileage > df.Mileage]
df = df[min_mileage < df.Mileage]
count_mileage = len(df.index)
print(count_price)
print("price removed: " + str(count_price - count_mileage) + " outliers")

# make year numeric, and remove outliers
maximum_age = 50
minimum_age = 0
df['Year'] = df['yearsold'] - df['Year']
df = df[maximum_age > df.Year]
df = df[minimum_age < df.Year]
count_age = len(df.index)
print(count_mileage)
print("price removed: " + str(count_mileage - count_age) + " outliers")

def try_to_find_from_engine(number, engine_value: str) -> str:
    if number != 0:
        return str(number).strip()
    found = re.search('v[0-9]+', engine_value, re.IGNORECASE)
    if found:
        return found.group()[1:].strip()
    else:
        return np.NaN

df['NumCylinders'] = df.apply(
```

```
        lambda row: try_to_find_from_engine(int(row['NumCylinders']), str(row[
    ])

engine_patterns = [
    (df['Engine'].str.contains('gas', case=False, regex=False, na=False),
    (df['Engine'].str.contains('diesel', case=False, regex=False, na=False)
]

criteria, values = zip(*engine_patterns)
df['FuelType'] = np.select(criteria, values, "unknown")

body_type_patterns = [
    (df['BodyType'].str.contains('hatchback', case=False, regex=False, na=
    (df['BodyType'].str.contains('coupe', case=False, regex=False, na=False)
    (df['BodyType'].str.contains('suv', case=False, regex=False, na=False)
    (df['BodyType'].str.contains('sedan', case=False, regex=False, na=False)
    (df['BodyType'].str.contains('pickup', case=False, regex=False, na=False)
    (df['BodyType'].str.contains('truck', case=False, regex=False, na=False)
    (df['BodyType'].str.contains('convertible', case=False, regex=False, n
    (df['BodyType'].str.contains('van', case=False, regex=False, na=False)
    (df['BodyType'].str.contains('wagon', case=False, regex=False, na=False)
]
criteria, values = zip(*body_type_patterns)
df['BodyType'] = np.select(criteria, values, "unknown")

drive_type_patterns = [
    (df['DriveType'].str.contains('RWD', case=False, regex=False, na=False)
    (df['DriveType'].str.contains('FWD', case=False, regex=False, na=False)
    (df['DriveType'].str.contains('4WD', case=False, regex=False, na=False)
    (df['DriveType'].str.contains('AWD', case=False, regex=False, na=False)
]

criteria, values = zip(*drive_type_patterns)
df['DriveType'] = np.select(criteria, values, "unknown")

def convert_engine_to_liters(engine_value: str):
    found = re.search('[0-9]\.[0-9]', engine_value, re.IGNORECASE)
    if found:
        return float(found.group())
    else:
        return np.NaN

df['Engine'] = df.apply(
    lambda row: convert_engine_to_liters(str(row['Engine'])), axis=1
)

df.dropna(subset=['NumCylinders', 'Engine'], inplace=True, axis=0)
```

```
count_engine_numCylinders = len(df.index)
print(count_age)
print("price removed: " + str(count_age - count_engine_numCylinders) + "

# get rid of rare models

# def delete_occurrences_fewer_than(data, model, threshold: int = 300) ->
#     if len(data[data['Model'] == model]) < threshold:
#         return 'Rare'
#     return model
#
#
# df['Model'] = df['Model'].apply(
#     lambda x: delete_occurrences_fewer_than(data=df, model=x)
# )
#df = df.groupby('Model').map((lambda x: x.value_counts() < 200), 'Other')

df = df.groupby('Make').filter(lambda x: len(x) > 100)
df['Model'] = df['Model'].mask(df['Model'].map(df['Model'].value_counts()))
df.drop('Make', inplace=True, axis=1)
df = df.groupby('NumCylinders').filter(lambda x: len(x) > 100)

df.dropna(subset=['Model'], inplace=True, axis=0)
#df.corr()

df = pd.get_dummies(data=df)
# for col in df.columns:
#     print(col)
print(df.shape)

# df_log2['pricesold'] = np.log2(df['pricesold'])

# splitting
# df_2019 = df[df["yearsold"] == 2019]
# df_2020 = df[df["yearsold"] == 2019]
# to_drop = ['yearsold']
# df_2019.drop(to_drop, inplace=True, axis=1)
# df_2020.drop(to_drop, inplace=True, axis=1)

year_sold_patterns = [
    (df['yearsold'] == 2018, 1),
    (df['yearsold'] == 2019, 2),
    (df['yearsold'] == 2020, 3),
]
```



```
criteria, values = zip(*year_sold_patterns)
df['yearsold'] = np.select(criteria, values, np.NAN)
df.dropna()

# df_2019_target = df_2019.pop('pricesold')
# df_2020_target = df_2020.pop('pricesold')

percentiles_price = np.percentile(df.pricesold, [25, 50, 75])

df_target = df.pop('pricesold')

# price is log2, shown by graph
df_target_log2 = np.log2(df_target) * 100
sns.displot(data=df_target_log2, kde=True)
plt.show()

# test data

x_train, x_test, y_train, y_test = train_test_split(df, df_target, test_si

# combine with price for test
x_test["pricesold"] = y_test

# make array from percentiles, to try
percentiles_array = [x_test[x_test["pricesold"].between(0, percentiles_pri
                      x_test[x_test["pricesold"].between(percentiles_price[
                      x_test[x_test["pricesold"].between(percentiles_price[
                      x_test[x_test["pricesold"].between(percentiles_price[

x_test.pop('pricesold')

# create separate array with price
percentiles_targets = []
for p in percentiles_array:
    print(len(p))
    percentiles_targets.append(p.pop('pricesold'))

y_train_log2 = np.log2(y_train) * 100

#train, val = train_test_split(train, test_size=0.2)
print(len(x_train), 'train examples')
#print(len(val), 'validation examples')
print(len(x_test), 'test examples')
```

```
# sklearn regression

# Defining MAPE function
def MAPE(Y_actual, Y_Predicted):
    mape = np.mean(np.abs((Y_actual - Y_Predicted) / Y_actual)) * 100
    return mape

# Defining prediction plot
def PLOT_ACTUAL_VS_PREDICTED(Y_actual, Y_predicted, line_of_best_fit=True)

    # add points to plot
    plt.scatter(Y_actual, Y_predicted)
    plt.xlabel('actual price')
    plt.ylabel('predicted price')
    plt.title('Best fit line with actual and predicted price')
    plt.legend()
    plt.tight_layout()

    if(line_of_best_fit):
        # find line of best fit
        a, b = np.polyfit(Y_actual.flatten(), Y_predicted.flatten(), 1)

        # add line of best fit to plot
        plt.plot(Y_actual, a * Y_actual + b, color='green', linewidth=6)
        plt.plot(Y_actual, Y_actual, color='orange', linewidth=6)
    plt.show()

# linear
lr = LinearRegression()
lr.fit(x_train, y_train_log2)
print("Linear Regression:")

# # predictions
# for d in percentiles_array:
#     pred_train_lr = lr.predict(x_train)
#     pred_train_lr = np.exp2(pred_train_lr * 0.01)
#     print(MAPE(y_train, pred_train_lr))

# predictions
pred_train_lr = lr.predict(x_train)
pred_train_lr = np.exp2(pred_train_lr * 0.01)
print("linear MAPE: "+str(MAPE(y_train, pred_train_lr)))
print("linear RSME: "+str(np.sqrt(mean_squared_error(y_train, pred_train_l
print("linear R2: "+str(r2_score(y_train, pred_train_lr)))

PLOT_ACTUAL_VS_PREDICTED(y_train.values, pred_train_lr)
```

```
pred_test_lr = lr.predict(x_test)
pred_test_lr = np.exp2(pred_test_lr * 0.01)
print("linear MAPE: "+str(MAPE(y_test, pred_test_lr)))
print("linear RSME: "+str(np.sqrt(mean_squared_error(y_test, pred_test_lr)))
print("linear R2: "+str(r2_score(y_test, pred_test_lr)))

PLOT_ACTUAL_VS_PREDICTED(y_test.values, pred_test_lr)

# ridge
rr = Ridge(alpha=0.01)
rr.fit(x_train, y_train_log2)
print("Ridge Regression:")

# predictions
pred_train_rr = rr.predict(x_train)
pred_train_rr = np.exp2(pred_train_rr * 0.01)
print("ridge MAPE: "+str(MAPE(y_train, pred_train_rr)))
print("ridge RSME: "+str(np.sqrt(mean_squared_error(y_train, pred_train_rr)))
print("ridge R2: "+str(r2_score(y_train, pred_train_rr)))

PLOT_ACTUAL_VS_PREDICTED(y_train.values, pred_train_rr)

pred_test_rr = rr.predict(x_test)
pred_test_rr = np.exp2(pred_test_rr * 0.01)
print("ridge MAPE: "+str(MAPE(y_test, pred_test_rr)))
print("ridge RSME: "+str(np.sqrt(mean_squared_error(y_test, pred_test_rr)))
print("ridge R2: "+str(r2_score(y_test, pred_test_rr)))

PLOT_ACTUAL_VS_PREDICTED(y_test.values, pred_test_rr)

# Lasso Regression
model_lasso = Lasso(alpha=0.01)
model_lasso.fit(x_train, y_train_log2)

# predictions
pred_train_lasso = model_lasso.predict(x_train)
pred_train_lasso = np.exp2(pred_train_lasso * 0.01)
print("lasso MAPE: "+str(MAPE(y_train, pred_train_lasso)))
print("lasso RSME: "+str(np.sqrt(mean_squared_error(y_train, pred_train_la
print("lasso R2: "+str(r2_score(y_train, pred_train_lasso)))

PLOT_ACTUAL_VS_PREDICTED(y_train.values, pred_train_lasso)

pred_test_lasso = model_lasso.predict(x_test)
pred_test_lasso = np.exp2(pred_test_lasso * 0.01)
print("lasso MAPE: "+str(MAPE(y_test, pred_test_lasso)))
```

```

print("lasso RSME: "+str(np.sqrt(mean_squared_error(y_test, pred_test_lasso)))
print("lasso R2: "+str(r2_score(y_test, pred_test_lasso)))

PLOT_ACTUAL_VS_PREDICTED(y_test.values, pred_test_lasso)

# Random Forest Regression
# important to test many times because it is stochastic
# get a list of models to evaluate
# def get_models():
#     models = dict()
#     # explore number of features from 1 to 7
#     for i in range(4, 6):
#         models[str(i)] = RandomForestClassifier(max_features=i)
#     return models
#
# # evaluate a given model using cross-validation
# def evaluate_model(model, X, y):
#     # define the evaluation procedure
#     cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=
#     # evaluate the model and collect the results
#     scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv, n_j
#     return scores

# Random Forest
rf = RandomForestRegressor(random_state=0)
rf.fit(x_train, y_train_log2)
print("RF Regression:")

# predictions
pred_train_rf = rf.predict(x_train)
pred_train_rf = np.exp2(pred_train_rf * 0.01)
print("RF MAPE: "+str(MAPE(y_train, pred_train_rf)))
print("RF RSME: "+str(np.sqrt(mean_squared_error(y_train, pred_train_rf)))
print("RF R2: "+str(r2_score(y_train, pred_train_rf)))

PLOT_ACTUAL_VS_PREDICTED(y_train.values, pred_train_rf)

pred_test_rf = rf.predict(x_test)
pred_test_rf = np.exp2(pred_test_rf * 0.01)
print("RF MAPE: "+str(MAPE(y_test, pred_test_rf)))
print("RF RSME: "+str(np.sqrt(mean_squared_error(y_test, pred_test_rf)))
print("RF R2: "+str(r2_score(y_test, pred_test_rf)))

PLOT_ACTUAL_VS_PREDICTED(y_test.values, pred_test_rf)

# Get MAPE boxplot

```

```
def graph_MAPE_by_price(array_to_predict):

    # create plot
    fig, ax = plt.subplots()
    index = np.arange(4)
    bar_width = 0.15
    opacity = 1
    colors = ['r', 'b', 'g', 'y']

    for i in range(4):

        results = []

        # lr
        pred_test = lr.predict(array_to_predict[i])
        pred_test = np.exp2(pred_test * 0.01)
        results.append(MAPE(percentiles_targets[i], pred_test))

        # rr
        pred_test = rr.predict(array_to_predict[i])
        pred_test = np.exp2(pred_test * 0.01)
        results.append(MAPE(percentiles_targets[i], pred_test))

        # lasso_r
        pred_test = model_lasso.predict(array_to_predict[i])
        pred_test = np.exp2(pred_test * 0.01)
        results.append(MAPE(percentiles_targets[i], pred_test))

        # rf
        pred_test = rf.predict(array_to_predict[i])
        pred_test = np.exp2(pred_test * 0.01)
        results.append(MAPE(percentiles_targets[i], pred_test))

        rects = plt.bar(index + bar_width * i, results, bar_width,
                        alpha=opacity,
                        color=colors[i],
                        label=str(i * 25) + 'th to ' + str((i+1) * 25) + 'th')

    plt.xlabel('Algorithm')
    plt.ylabel('MAPE')
    plt.title('MAPE by Algorithm for price percentile')
    plt.xticks(index + bar_width, ('Linear', 'Ridge', 'Lasso', 'RandomFore
plt.legend(bbox_to_anchor=(1.1, 1.05))
plt.tight_layout()
plt.show()
```

```
graph_MAPE_by_price(percentiles_array)
```

```
# numeric features
# numeric_feature_names = ['mileage', 'hp', 'year']
# numeric_features = df[numeric_feature_names]
# numeric_features.sort_values(['hp'], axis=0, ascending=True, inplace=True)
# print(numeric_features)
# print(numeric_features.dtypes)
# print(numeric_features.shape)
```

```
# plot the density histograms
```

```
def increment_year(dataframe):
    dataframe["Year"] += 1
    dataframe["yearsold"] += 1
    dataframe["Mileage"] += 14263
```

```
# return average prediction list from dataframe array
# feed for one year
# return list for 2019, 2018
def get_average_prediction(df_arr):
```

```
    def percentage_difference(result_old, result_new):
        result_old = np.average(np.exp2(result_old * 0.01))
        result_new = np.average(np.exp2(result_new * 0.01))
        decrease = result_old - result_new
        return (decrease / result_old) * 100
```

```
    linreg = percentage_difference(lr.predict(df_arr[0]), lr.predict(df_ar
    ridgereg = percentage_difference(rr.predict(df_arr[0]), rr.predict(df_
    lassoreg = percentage_difference(model_lasso.predict(df_arr[0]), model
    randomreg = percentage_difference(rf.predict(df_arr[0]), rf.predict(df_
```

```
    # return tuple with all data for one year
    # actual data too?
    return linreg, ridgereg, lassoreg, randomreg
```

```
# 2019 dataset
df_2019 = df[df["yearsold"] == 2]

# split into 2-4
df_2019_newer = df_2019[df_2019["Year"].between(2, 4)]
```

```
df_2020_newer = df_2019_newer.copy()
increment_year(df_2020_newer)

newer_depreciation = get_average_prediction([df_2019_newer, df_2020_newer])

# split into 13-15
df_2019_older = df_2019[df_2019["Year"].between(13, 15)]

df_2020_older = df_2019_older.copy()
increment_year(df_2020_older)

older_depreciation = get_average_prediction([df_2019_older, df_2020_older])

# create plot
fig, ax = plt.subplots()
index = np.arange(4)
bar_width = 0.35
opacity = 0.8

rects1 = plt.bar(index, newer_depreciation, bar_width,
                  alpha=opacity,
                  color='b',
                  label='2-4 year old cars')

rects2 = plt.bar(index + bar_width, older_depreciation, bar_width,
                  alpha=opacity,
                  color='g',
                  label='13-15 year old cars')

plt.xlabel('Algorithm')
plt.ylabel('Depreciation')
plt.title('Depreciation by Algorithm')
plt.xticks(index + bar_width, ('Linear', 'Ridge', 'Lasso', 'RandomForest'))
plt.legend()
plt.tight_layout()
plt.show()

# find the most common occurrences all columns in dataframe
# most common mileage: 100,000
sns.distplot(df['Mileage'], hist=True, kde=True, kde_kws={'linewidth': 3},

# most common Year: 5-15
sns.distplot(df['Year'], hist=True, kde=True, kde_kws={'linewidth': 3}, la
plt.show()

# most common Engine: 3.8 - 4.8
sns.distplot(df['Engine'], hist=True, kde=True, kde_kws={'linewidth': 3},
```

```

print(list(df))

model_results = [(pred_test_rr, 'RidgeRegression'), (pred_test_lasso, 'Las

for i in range(4):
    # Draw the density plot
    result = model_results[i]
    sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1}

sns.distplot(y_test.values, hist=False, kde=True, kde_kws={'linewidth': 1}

# Plot formatting
plt.legend(title='Price')
plt.title('Density Plot with Multiple Model Prediction Prices')
plt.xlabel('Price (USD)')
plt.ylabel('Density')
plt.show()


```

```

pricesold      int64
yearsold       int64
Mileage        int64
Make           object
Model          object
Year           int64
Engine         object
BodyType       object
NumCylinders   int64
DriveType      object
dtype: object
122144
price removed: 11743 outliers
110401
price removed: 13950 outliers
96451
price removed: 15902 outliers
80549
price removed: 38562 outliers
.....

```


(26422, 129)



1360

1320

1277

1377

21137 train examples

5285 test examples

WARNING:matplotlib.legend:No artists with labels found to put in legend. 1


Linear Regression:

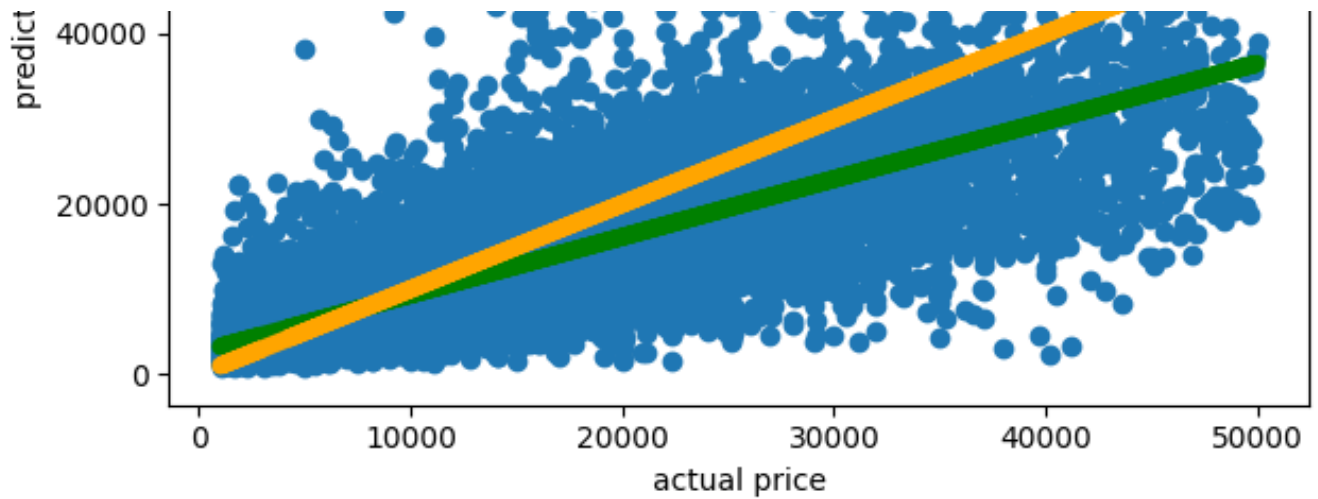
linear MAPE: 44.890548775416754

linear RSME: 5820.681978626334

linear R2: 0.6457558729716408


Best fit line with actual and predicted price





```
WARNING:matplotlib.legend:No artists with labels found to put in legend.
linear MAPE: 46.325834113513146
linear RSME: 5969.136375597109
linear R2: 0.6499672700215979
```


Best fit line with actual and predicted price



```
WARNING:matplotlib.legend:No artists with labels found to put in legend.
Ridge Regression:
ridge MAPE: 44.89056805904398
ridge RSME: 5820.616925861155
ridge R2: 0.6457637910923425
```


Best fit line with actual and predicted price





```
WARNING:matplotlib.legend:No artists with labels found to put in legend.
ridge MAPE: 46.325875278531505
ridge RSME: 5969.059965121119
ridge R2: 0.6499762314507358
```


Best fit line with actual and predicted price



actual price


```
WARNING:matplotlib.legend:No artists with labels found to put in legend.
lasso MAPE: 44.91314130816676
lasso RSME: 5816.5046086727
lasso R2: 0.6462641562846697
```


Best fit line with actual and predicted price



```
WARNING:matplotlib.legend:No artists with labels found to put in legend.
lasso MAPE: 46.36972694931304
lasso RSME: 5966.075482583576
lasso R2: 0.6503261621518037
```

Best fit line with actual and predicted price





RF Regression:


WARNING:matplotlib.legend:No artists with labels found to put in legend. 1

RF MAPE: 12.869616615915424

RF RSME: 1982.8458212550986

RF R2: 0.9588913884761017

Best fit line with actual and predicted price



WARNING:matplotlib.legend:No artists with labels found to put in legend. 1




RF MAPE: 38.81974326298702

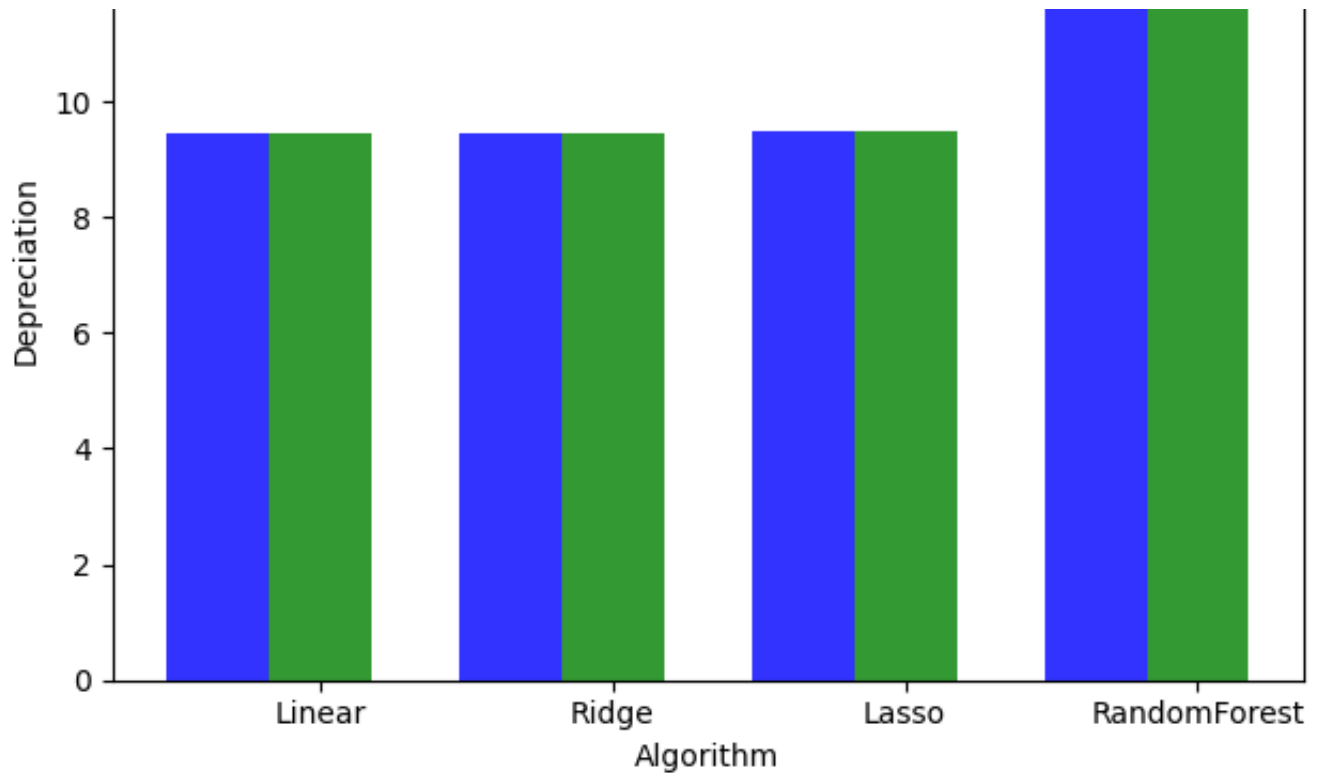
RF RSME: 4760.2520313657415

RF R2: 0.7773895504128551

Best fit line with actual and predicted price







<ipython-input-37-3e5cefc59bde>:480: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Mileage'], hist=True, kde=True, kde_kws={'linewidth': 3})
```

<ipython-input-37-3e5cefc59bde>:483: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Year'], hist=True, kde=True, kde_kws={'linewidth': 3}, 1
```





<ipython-input-37-3e5cefc59bde>:487: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Engine'], hist=True, kde=True, kde_kws={'linewidth': 3},
['yearsold', 'Mileage', 'Year', 'Engine', 'Model_1500', 'Model_3-Series', '

```

<ipython-input-37-3e5cefc59bde>:497: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1}, 1

```

<ipython-input-37-3e5cefc59bde>:497: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1}, 1

```

<ipython-input-37-3e5cefc59bde>:497: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with

similar flexibility) or `kdeplot` (an axes-level function for kernel density

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1}, 1
<ipython-input-37-3e5cefc59bde>:497: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1}, 1
<ipython-input-37-3e5cefc59bde>:499: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(y_test.values, hist=False, kde=True, kde_kws={'linewidth': 1
```



FILE (USD)

```
sns.distplot(df['Mileage'], hist=True, kde=True, kde_kws={'linewidth': 3}, lab  
plt.show())
```

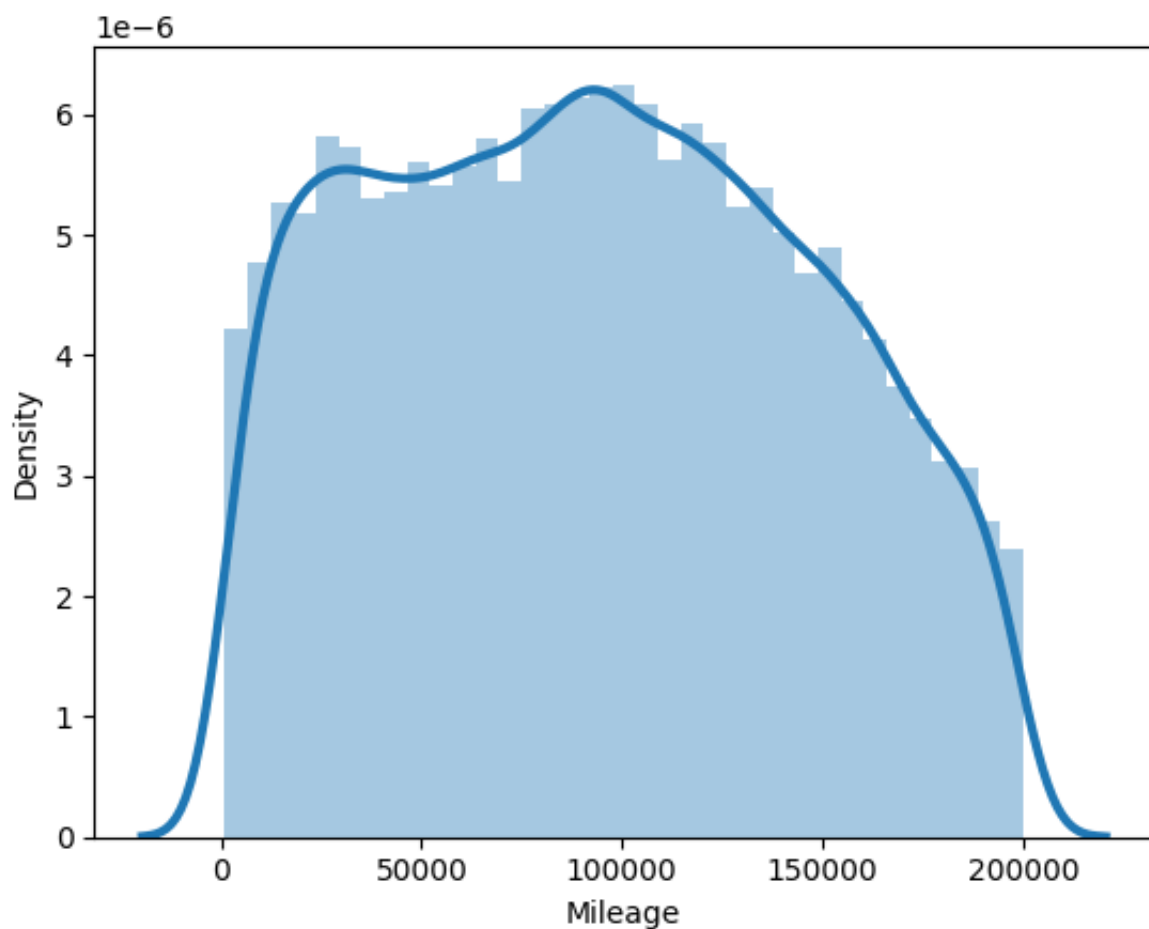
<ipython-input-16-b416db703a07>:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Mileage'], hist=True, kde=True, kde_kws={'linewidth': 3})
```



```
sns.distplot(df['Year'], hist=True, kde=True, kde_kws={'linewidth': 3}, label=
plt.show())
```


<ipython-input-17-f10485ea5b97>:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Year'], hist=True, kde=True, kde_kws={'linewidth': 3}, 1
```



```
sns.distplot(df['Engine'], hist=True, kde=True, kde_kws={'linewidth': 3}, label=
plt.show())
```


```
<ipython-input-18-0da71dec8133>:1: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Engine'], hist=True, kde=True, kde_kws={'linewidth': 3},
```



```
import seaborn as sns
import matplotlib.pyplot as plt
```

```
# Most common mileage: 100,000
```

```
sns.displot(df['Mileage'], kde=True, label="Mileage")
```

```
sns.kdeplot(df['Mileage'], linewidth=3) # Specify the linewidth directly for
```

```
plt.show()

# Most common Year: 5-15
sns.displot(df['Year'], kde=True, label="Year")
sns.kdeplot(df['Year'], linewidth=3)
plt.show()


# Most common Engine: 3.8 - 4.8
sns.displot(df['Engine'], kde=True, label="Engine")
sns.kdeplot(df['Engine'], linewidth=3)
plt.show()



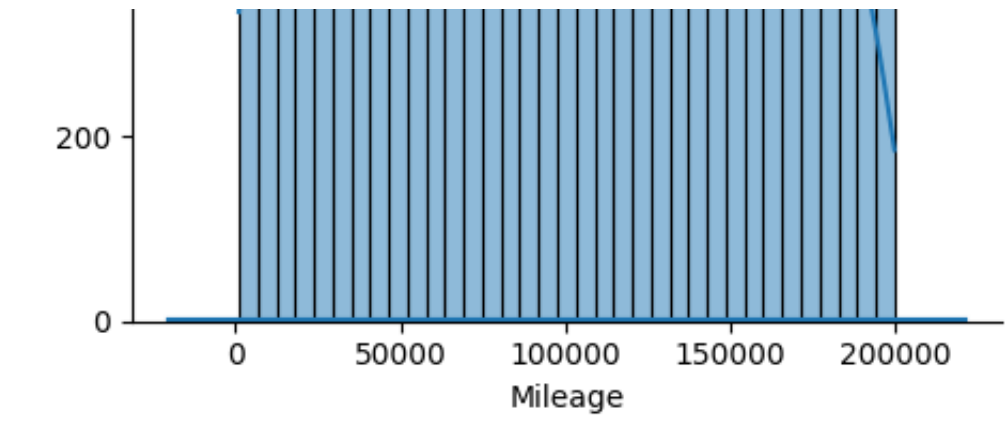
print(list(df))


model_results = [
    (pred_test_rr, 'RidgeRegression'),
    (pred_test_lasso, 'LassoRegression'),
    (pred_test_rf, 'RandomForest'),
    (pred_test_lr, 'LinearRegression')
]




for i in range(4):
    # Draw the density plot
    result = model_results[i]
    sns.displot(result[0], kde=True, label=result[1])
    sns.kdeplot(result[0], linewidth=1)
    plt.show()



sns.displot(y_test.values, kde=True, label='Actual')
sns.kdeplot(y_test.values, linewidth=1)
plt.show()
```











```

for i in range(4):
    result = model_results[i]
    sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1}, la
sns.distplot(y_test.values, hist=False, kde=True, kde_kws={'linewidth': 1}, la
plt.legend(title='Price')
plt.title('Multiple Model Prediction Prices')
plt.xlabel('Price (USD)')
plt.ylabel('Density')
plt.show()

```

<ipython-input-30-09c6eb27e06f>:3: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density estimation).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```

sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1}, l
<ipython-input-30-09c6eb27e06f>:3: UserWarning:

```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density estimation).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```

sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1}, l
<ipython-input-30-09c6eb27e06f>:3: UserWarning:

```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density estimation).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```

sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1}, l
<ipython-input-30-09c6eb27e06f>:3: UserWarning:

```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density estimation).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

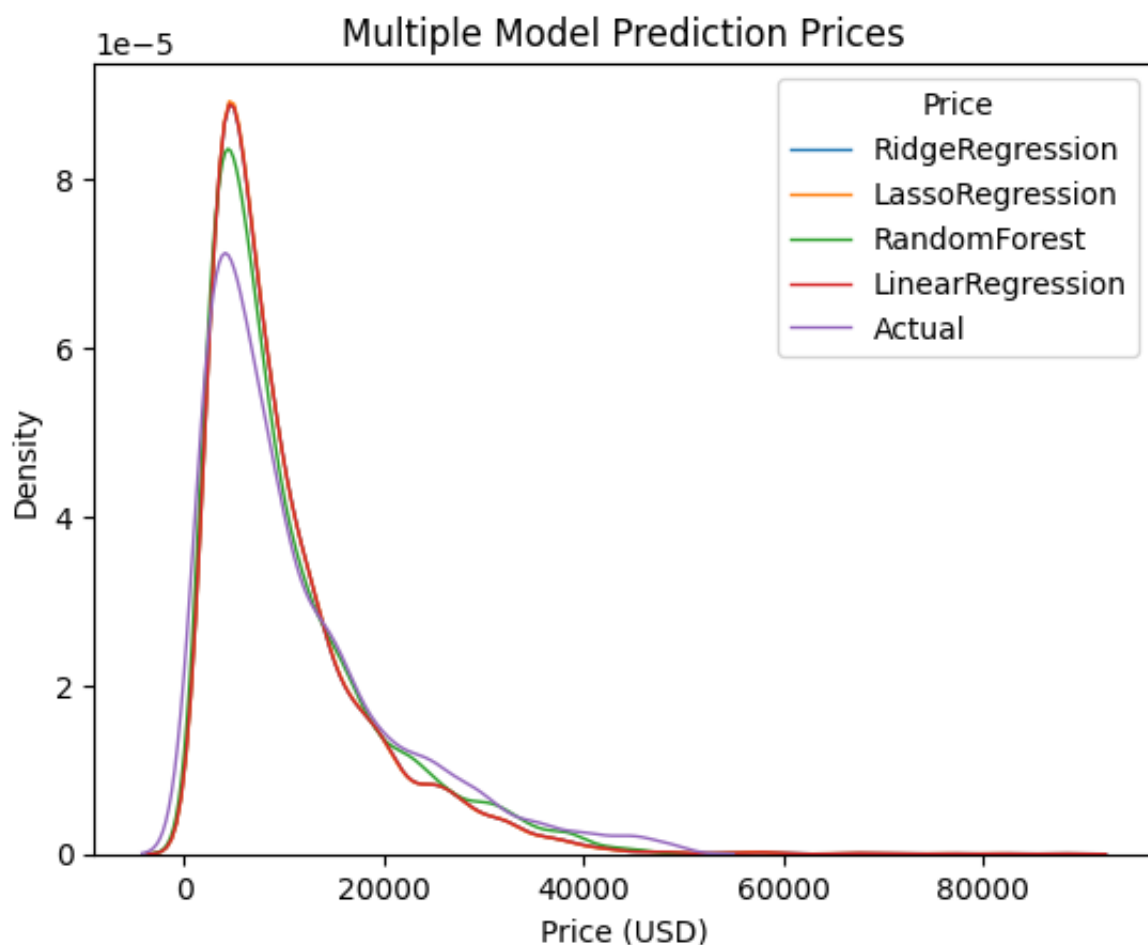
```
sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1}, 1  
<ipython-input-30-09c6eb27e06f>:4: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(y_test.values, hist=False, kde=True, kde_kws={'linewidth': 1
```




✓ Customizing Charts

```
# Plot
rects1 = plt.bar(index, newer_depreciation, bar_width, alpha=opacity, color='orange')
rects2 = plt.bar(index + bar_width, older_depreciation, bar_width, alpha=opacity, color='blue')

# Customizing plot
plt.xlabel('Algorithms')
plt.ylabel('Depreciation')
plt.title('Depreciation by different Algorithms')
plt.xticks(index + bar_width, ('Linear', 'Ridge', 'Lasso', 'RandomForest'))
plt.legend()
plt.tight_layout()

# Show plot
plt.show()
```



```
# Define colors
colors = ['orange', 'red', 'green', 'black', 'purple']

# Plot KDE plots for each model result
for i in range(4):
    result = model_results[i]
```

```

sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1, 'co

# Plot KDE plot for actual values
sns.distplot(y_test.values, hist=False, kde=True, kde_kws={'linewidth': 1, 'co

# Customizing plot
plt.legend(title='Price')
plt.title('Multiple Model Prediction Prices')
plt.xlabel('Price')
plt.ylabel('Density')
plt.show()

```

<ipython-input-43-e71a4df6781a>:7: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density estimation).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```

sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1, 'co
<ipython-input-43-e71a4df6781a>:7: UserWarning:

```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density estimation).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```

sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1, 'co
<ipython-input-43-e71a4df6781a>:7: UserWarning:

```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density estimation).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```

sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1, 'co
<ipython-input-43-e71a4df6781a>:7: UserWarning:

```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density estimation).

For a guide to updating your code to use the new functions, please see

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(result[0], hist=False, kde=True, kde_kws={'linewidth': 1, 'c  
<ipython-input-43-e71a4df6781a>:10: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(y_test.values, hist=False, kde=True, kde_kws={'linewidth': 1
```




```
import matplotlib.pyplot as plt
import numpy as np

def graph_MAPE_by_price(array_to_predict):

    # Create plot
    fig, ax = plt.subplots()
    index = np.arange(4)
    bar_width = 0.15
    opacity = 1
    colors = ['orange', 'red', 'blue', 'purple'] # Modify colors here

    for i in range(4):
        results = []

        # lr
        pred_test = lr.predict(array_to_predict[i])
        pred_test = np.exp2(pred_test * 0.01)
        results.append(MAPE(percentiles_targets[i], pred_test))

        # rr
        pred_test = rr.predict(array_to_predict[i])
        pred_test = np.exp2(pred_test * 0.01)
        results.append(MAPE(percentiles_targets[i], pred_test))

        # lasso_r
        pred_test = model_lasso.predict(array_to_predict[i])
        pred_test = np.exp2(pred_test * 0.01)
        results.append(MAPE(percentiles_targets[i], pred_test))


        # rf
        pred_test = rf.predict(array_to_predict[i])
        pred_test = np.exp2(pred_test * 0.01)
        results.append(MAPE(percentiles_targets[i], pred_test))

        rects = plt.bar(index + bar_width * i, results, bar_width,
                        alpha=opacity,
                        color=colors[i],
                        label=str(i * 25) + 'th to ' + str((i+1) * 25) + 'th per

    plt.xlabel('Algorithms')
    plt.ylabel('MAPE')
    plt.title('Price percentile by MAPE Algorithm')
```

```
plt.xticks(index + bar_width, ('Linear', 'Ridge', 'Lasso', 'RandomForest'))  
plt.legend(bbox_to_anchor=(1.1, 1.05))  
plt.tight_layout()  
plt.show()
```

```
graph_MAPE_by_price(percentiles_array)
```



```
# Defining prediction plot
def PLOT_ACTUAL_VS_PREDICTED(Y_actual, Y_predicted, line_of_best_fit=True):

    # Add points to plot with orange color
    plt.scatter(Y_actual, Y_predicted, color='orange', label='Predicted vs Act

    plt.xlabel('actual price')
    plt.ylabel('predicted price')
    plt.title('Best fit line with actual and predicted price')
    plt.legend()
    plt.tight_layout()

    if(line_of_best_fit):
        # find line of best fit
        a, b = np.polyfit(Y_actual.flatten(), Y_predicted.flatten(), 1)

        # add line of best fit to plot
        plt.plot(Y_actual, a * Y_actual + b, color='green', linewidth=6)
        plt.plot(Y_actual, Y_actual, color='red', linewidth=6)
    plt.show()

# Fit your regression models (lr, rr, model_lasso) before running this code

# Extract feature names
feature_names = df.columns

# Create a table header
print("MODEL\tFEATURE COEFFICIENTS")

# Loop through each model
for model_name, model in [("LINEAR REGRESSION", lr),
                           ("RIDGE REGRESSION", rr),
                           ("LASSO REGRESSION", model_lasso)]:

    # Get coefficients from the model
    coefficients = model.coef_

    # Print model name
    print(model_name)

    # Print feature coefficients
    for feature_name, coef in zip(feature_names, coefficients):
        print(f"{feature_name}\t{coef:.4f}")

    print("\n") # Add a newline between models
```

```

MODEL    FEATURE COEFFICIENTS
LINEAR REGRESSION
yearsold      9.3502
Mileage -0.0012
Year    -5.5789
Engine  29.6578
Model_1500      8993526.0152
Model_3-Series  8993571.6430
Model_300 Series      8993526.8578
Model_300-Series      8993633.2913
Model_4Runner  8993605.4687
Model_5-Series  8993567.2221
Model_6-Series  8993619.4876
Model_7-Series  8993566.7584
Model_911      8993793.5632
Model_A4       8993510.1030
Model_A6       8993553.1100
Model_Acadia   8993552.9886
Model_Accord   8993533.1636
Model_Altime   8993484.0556
Model_Blazer   8993538.4395
Model_Boxster  8993611.3963
Model_Bronco   8993626.1665
Model_C-10     8993663.4768
Model_C-Class  8993540.2520
Model_C/K Pickup 1500 8993561.8117
Model_CR-V     8993538.2315
Model_CTS      8993543.5640
Model_Camaro   8993554.1613
Model_Camry    8993538.0222
Model_Cayenne  8993608.3271
Model_Challenger 8993553.1737
Model_Charger  8993526.2959
Model_Cherokee 8993548.6398
Model_Civic    8993519.1144
Model_Corolla  8993516.9811
Model_Corvette 8993578.0350
Model_Crown Victoria 8993426.5565
Model_DeVille  8993469.6129
Model_Durango  8993500.4397
Model_E-Class  8993570.1419
Model_E-Series Van 8993501.5954
Model_Edge     8993546.3422
Model_Eldorado 8993530.9130
Model_Enclave  8993560.1990
Model_Equinox  8993501.3459
Model_Escalade 8993580.1012
Model_Escape   8993491.2169
Model_Excursion 8993595.4337
Model_Expedition 8993522.7769
Model_Explorer 8993522.1694
Model_Express   8993507.0000
Model_F-150     8993552.4776
Model_F-250     8993523.6694

```

```

Model_F-350      8993539.3002
Model_Firebird   8993576.3094
Model_Focus      8993451.6029
Model_Forester   8993495.6941
Model_Fusion     8993504.6685
Model_Grand_Caravan 8993520.0000

```

```

# Define the headers for the table
print("PERCENTILE\tSTART PRICE\tEND PRICE")

# Define the percentile labels and corresponding start and end prices
percentile_labels = ["0TH-25TH", "25TH-50TH", "50TH-75TH", "75TH-100TH"]
start_prices = [min_price, percentiles_price[0], percentiles_price[1], percentiles_price[2]]
end_prices = [percentiles_price[0], percentiles_price[1], percentiles_price[2], percentiles_price[3]]

# Iterate through the percentiles and print the corresponding start and end prices
for label, start_price, end_price in zip(percentile_labels, start_prices, end_prices):
    print(f"{label}\t\t\t${start_price:.2f}\t\t\t${end_price:.2f}")

```

PERCENTILE	START PRICE	END PRICE
0TH-25TH	\$1000.00	\$4050.00
25TH-50TH	\$4050.00	\$8050.00
50TH-75TH	\$8050.00	\$15500.00
75TH-100TH	\$15500.00	\$50000.00

```
df.describe()
```

	yearsold	Mileage	Year	Engine	Model_1500	Model_1500
count	26422.000000	26422.000000	26422.000000	26422.000000	26422.000000	26422.000000
mean	2.372228	93189.599198	14.163311	4.319590	0.004352	0.004352
std	0.503428	52704.345670	9.778808	1.434038	0.065830	0.065830
min	1.000000	1011.000000	1.000000	1.000000	0.000000	0.000000
25%	2.000000	48851.000000	6.000000	3.400000	0.000000	0.000000
50%	2.000000	92000.000000	13.000000	4.300000	0.000000	0.000000
75%	3.000000	134898.250000	19.000000	5.500000	0.000000	0.000000
max	3.000000	199999.000000	49.000000	8.800000	1.000000	1.000000

8 rows × 128 columns